

HW_4_2_GAN

March 12, 2025

1 Understanding GANs

Imagine walking into an art studio where something fascinating is happening: there's an artist and an art critic engaged in an unusual competition. The artist is trying to create paintings that look exactly like works from a famous gallery, while the critic's job is to spot the difference between real gallery paintings and the artist's creations. This back-and-forth between artist and critic perfectly captures how GANs (Generative Adversarial Networks) work.

1.1 Core Intuition

In our Swiss Roll example, we can think of it this way: The artist (Generator) is trying to create points that form the distinctive spiral pattern of a Swiss Roll, while the critic (Discriminator) needs to distinguish between points from the real Swiss Roll distribution and the artist's generated points. Through this competition, both networks improve – the artist gets better at creating realistic points, and the critic becomes more discerning.

1.2 The Mathematical Dance

This artistic competition translates into a mathematical game. The Generator (G) and Discriminator (D) play what we call a minimax game, represented by this equation:

$$\min_G \max_D \mathbb{E}_x[\log D(x)] + \mathbb{E}_z[\log(1 - D(G(z)))]$$

Breaking this down in simple terms:

- $D(x)$ is the critic's judgment of real data (should be close to 1)
- $G(z)$ is the artist creating new data from random noise z
- $D(G(z))$ is the critic's judgment of generated data (should be close to 0)
- The Generator tries to minimize this difference
- The Discriminator tries to maximize it

1.2.1 The Latent Space (potential homework? to explore)

The noise input to our Generator isn't just random numbers – it's a structured space where:

- Similar points generate similar outputs
- We can interpolate between points
- Different dimensions might control different features

1.3 Key Papers

For a deeper dive into GANs, consider reading:

- “Generative Adversarial Networks” by Goodfellow et al. (2014)
- “Wasserstein GAN” by Arjovsky et al. (2017)

```
[2]: import torch
import torch.nn as nn
import numpy as np
import matplotlib.pyplot as plt
from tqdm import tqdm

# Set random seed for reproducibility
torch.manual_seed(42)
np.random.seed(42)

# Generate Swiss Roll data (same as before)
def generate_swiss_roll(n_samples=1000):
    """
    Generate 2D Swiss Roll data points
    Args:
        n_samples: Number of points to generate
    Returns:
        torch.Tensor of shape (n_samples, 2)
    """
    t = 1.5 * np.pi * (1 + 2 * torch.rand(n_samples))
    x = t * torch.cos(t)
    y = t * torch.sin(t)
    data = torch.stack([x, y], dim=1) / 15.0 # Scale down the data
    return data

data = generate_swiss_roll(n_samples=10000)
plt.scatter(data[:, 0], data[:, 1], alpha=0.5, s=10)
plt.title("Sythetic Swiss Roll Distribution")

class Generator(nn.Module):
    """
    Generator network that transforms random noise into 2D points
    to mimic the Swiss Roll distribution
    """
    def __init__(self, latent_dim=2, hidden_dim=256):
        super().__init__()
        self.net = nn.Sequential(
```

```

        nn.Linear(latent_dim, hidden_dim),
        nn.LeakyReLU(0.2),
        nn.Linear(hidden_dim, hidden_dim),
        nn.LeakyReLU(0.2),
        nn.Linear(hidden_dim, hidden_dim),
        nn.LeakyReLU(0.2),
        nn.Linear(hidden_dim, 2), # Output 2D coordinates
        nn.Tanh(), # Bound outputs to [-1, 1]
    )

    def forward(self, z):
        return self.net(z)

class Discriminator(nn.Module):
    """
    Discriminator network that tries to distinguish real Swiss Roll points
    from generated ones
    """

    def __init__(self, hidden_dim=256):
        super().__init__()
        self.net = nn.Sequential(
            nn.Linear(2, hidden_dim),
            nn.LeakyReLU(0.2),
            nn.Linear(hidden_dim, hidden_dim),
            nn.LeakyReLU(0.2),
            nn.Linear(hidden_dim, hidden_dim),
            nn.LeakyReLU(0.2),
            nn.Linear(hidden_dim, 1),
            nn.Sigmoid(), # Output probability between 0 and 1
        )

    def forward(self, x):
        return self.net(x)

def train_gan(n_steps=10000, batch_size=128, lr=2e-4, latent_dim=2):
    """
    Train the GAN on Swiss Roll data
    Args:
        n_steps: Number of training steps
        batch_size: Batch size for training
        lr: Learning rate
        latent_dim: Dimension of noise input to generator
    Returns:
        Trained generator and discriminator models, training losses, device
    """

```

```

"""
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

# Initialize networks and optimizers
generator = Generator(latent_dim=latent_dim).to(device)
discriminator = Discriminator().to(device)

# Use Adam optimizer with beta parameters recommended for GANs
g_optimizer = torch.optim.Adam(generator.parameters(), lr=lr, betas=(0.5, 0.
↪999))
d_optimizer = torch.optim.Adam(
    discriminator.parameters(), lr=lr, betas=(0.5, 0.999)
)

criterion = nn.BCELoss()

# Training loop
g_losses, d_losses = [], []
pbar = tqdm(range(n_steps), desc="Training GAN")

for step in pbar:
    # Train Discriminator
    for _ in range(1): # Can adjust D/G training ratio
        # Real data
        real_data = generate_swiss_roll(batch_size).to(device)
        real_labels = torch.ones(batch_size, 1).to(device)

        # Generated data
        z = torch.randn(batch_size, latent_dim).to(device)
        fake_data = generator(z).detach()
        fake_labels = torch.zeros(batch_size, 1).to(device)

        # Train on real and fake data
        d_optimizer.zero_grad()
        d_real_loss = criterion(discriminator(real_data), real_labels)
        d_fake_loss = criterion(discriminator(fake_data), fake_labels)
        d_loss = d_real_loss + d_fake_loss
        d_loss.backward()
        d_optimizer.step()

    # Train Generator
    for _ in range(1):
        z = torch.randn(batch_size, latent_dim).to(device)
        g_optimizer.zero_grad()
        fake_data = generator(z)
        g_loss = criterion(discriminator(fake_data), real_labels) # Try to
↪fool D

```

```

        g_loss.backward()
        g_optimizer.step()

        # Record losses
        g_losses.append(g_loss.item())
        d_losses.append(d_loss.item())

        if step % 100 == 0:
            pbar.set_postfix({"G_Loss": g_loss.item(), "D_Loss": d_loss.item()})

    return generator, discriminator, g_losses, d_losses, device

# Train the model
generator, discriminator, g_losses, d_losses, device = train_gan()

# Plot training losses
plt.plot(g_losses, label="Generator")
plt.plot(d_losses, label="Discriminator")
plt.xlabel("Steps")
plt.ylabel("Loss")
plt.legend()
plt.show()

def visualize_gan_samples(generator, n_samples=1000):
    """
    Visualize samples from the trained generator alongside real data
    Args:
        generator: Trained generator model
        n_samples: Number of points to generate
    """
    device = next(generator.parameters()).device

    # Generate samples
    with torch.no_grad():
        z = torch.randn(n_samples, 2).to(device)
        fake_data = generator(z).cpu().numpy()

    # Get real data for comparison
    real_data = generate_swiss_roll(n_samples).numpy()

    # Plot
    plt.figure(figsize=(12, 5))

    plt.subplot(1, 2, 1)

```

```

plt.scatter(real_data[:, 0], real_data[:, 1], c="blue", alpha=0.5,
↳label="Real")
plt.title("Real Swiss Roll")
plt.legend()
plt.axis("equal")

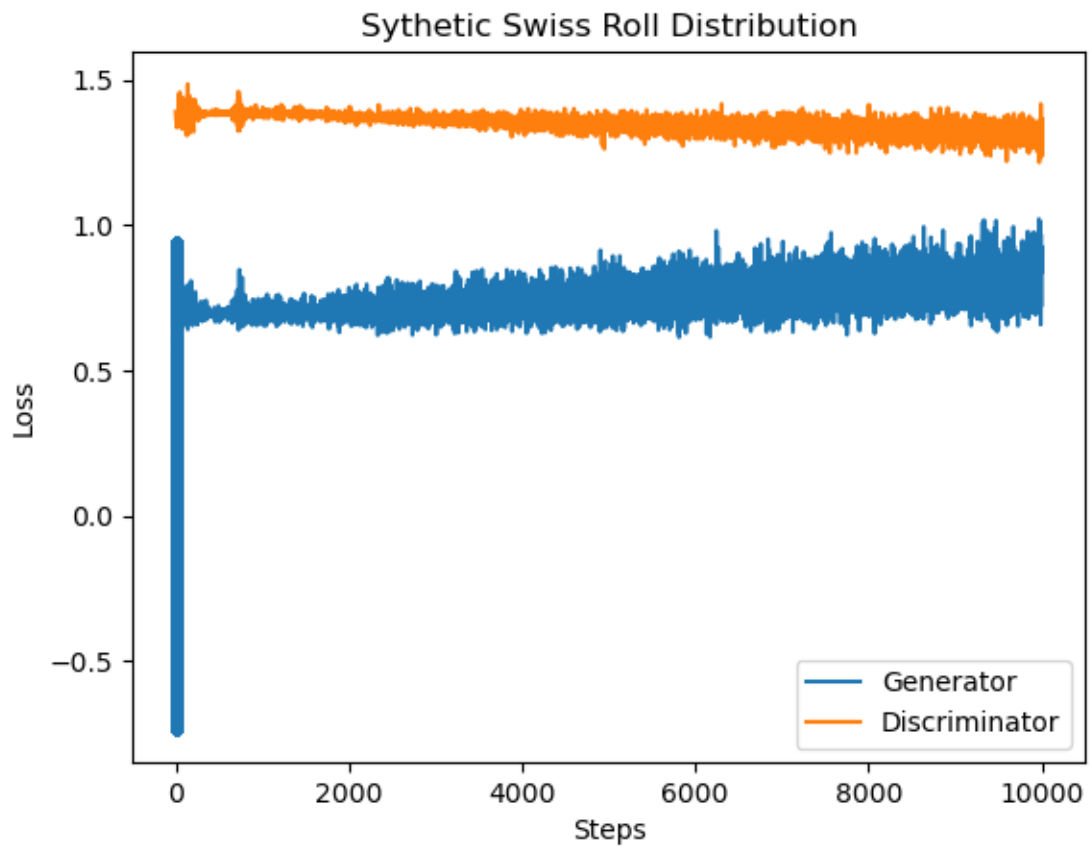
plt.subplot(1, 2, 2)
plt.scatter(fake_data[:, 0], fake_data[:, 1], c="red", alpha=0.5,
↳label="Generated")
plt.title("GAN Generated")
plt.legend()
plt.axis("equal")

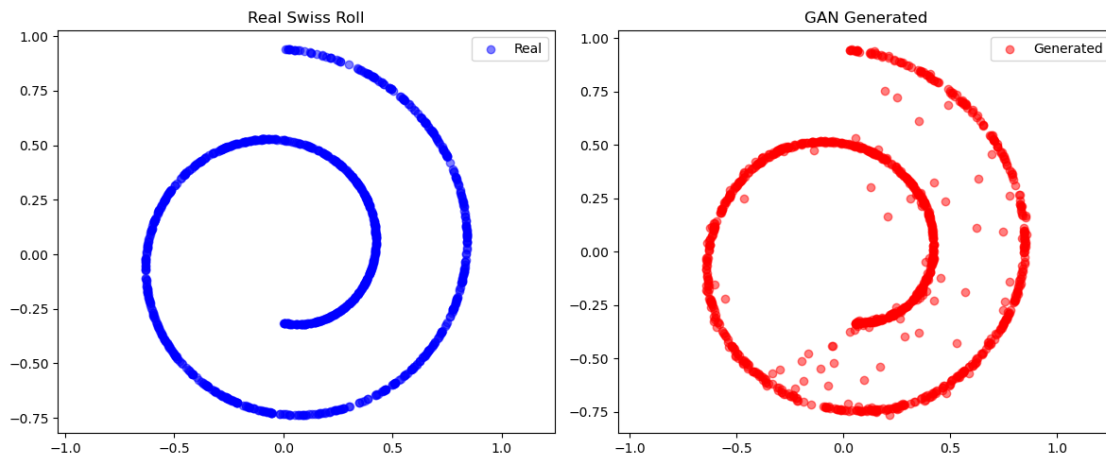
plt.tight_layout()
plt.show()

# Visualize results
visualize_gan_samples(generator)

```

Training GAN: 100%| | 10000/10000 [07:29<00:00, 22.25it/s,
G_Loss=0.817, D_Loss=1.29]





2 Generative Adversarial Network (GAN) for Swiss Roll Data

In this notebook you will implement—in part—the training loop for a GAN on a 2D Swiss Roll dataset.

Overview:

- **Generator:**
Transforms random noise into 2D points that mimic the Swiss Roll data.
- **Discriminator:**
Distinguishes real Swiss Roll points from those generated by the generator.
- **Training Objectives:**
 - **Discriminator:** Maximize the probability of assigning correct labels to real and fake samples.
 - **Generator:** Fool the discriminator by generating samples that are labeled as real.

Your Task:

Review and complete the GAN training loop. In particular, you will work on the part where: - Real and fake labels are created. - The discriminator loss is computed on real data (with true labels) and on fake data (with fake labels). - The generator loss is computed by trying to convince the discriminator that generated samples are real.

Study the provided hints and inline comments carefully.

```
[3]: import torch
import torch.nn as nn
import torch.optim as optim # Import the optimizer module
import numpy as np
import matplotlib.pyplot as plt
```

```

from tqdm import tqdm

# Set random seed for reproducibility.
torch.manual_seed(42)
np.random.seed(42)

```

2.1 Data Generation

We define a function to generate 2D Swiss Roll data points. This will serve as our real data distribution.

```


[4]: def generate_swiss_roll(n_samples=1000):
    """
    Generate 2D Swiss Roll data points.

    Args:
        n_samples: Number of points to generate.

    Returns:
        torch.Tensor of shape (n_samples, 2)
    """
    t = 1.5 * np.pi * (1 + 2 * torch.rand(n_samples))
    x = t * torch.cos(t)
    y = t * torch.sin(t)
    data = torch.stack([x, y], dim=1) / 15.0 # Scale down the data
    return data

# Visualize some Swiss Roll data.
data = generate_swiss_roll(n_samples=10000)
plt.scatter(data[:, 0].detach().numpy(), data[:, 1].detach().numpy(), alpha=0.
    ↪5, s=10)
plt.title("Synthetic Swiss Roll Distribution")
plt.show()

```

2.2 Network Architectures

Below are the Generator and Discriminator network definitions.

- **Generator:**
Takes a latent noise vector and outputs a 2D sample.
The final Tanh activation bounds outputs to $[-1, 1]$.
- **Discriminator:**
Takes a 2D input and outputs a scalar probability (after Sigmoid) indicating whether the input is real.

```
[5]: class Generator(nn.Module):
    """
    Generator network that transforms random noise into 2D points
    to mimic the Swiss Roll distribution.
    """

    def __init__(self, latent_dim=2, hidden_dim=256):
        super().__init__()
        self.net = nn.Sequential(
            nn.Linear(latent_dim, hidden_dim),
```

```

        nn.LeakyReLU(0.2),
        nn.Linear(hidden_dim, hidden_dim),
        nn.LeakyReLU(0.2),
        nn.Linear(hidden_dim, hidden_dim),
        nn.LeakyReLU(0.2),
        nn.Linear(hidden_dim, 2), # Output 2D coordinates
        nn.Tanh(), # Bound outputs to [-1, 1]
    )

    def forward(self, z):
        return self.net(z)

class Discriminator(nn.Module):
    """
    Discriminator network that tries to distinguish real Swiss Roll points
    from generated ones.
    """

    def __init__(self, hidden_dim=256):
        super().__init__()
        self.net = nn.Sequential(
            nn.Linear(2, hidden_dim),
            nn.LeakyReLU(0.2),
            nn.Linear(hidden_dim, hidden_dim),
            nn.LeakyReLU(0.2),
            nn.Linear(hidden_dim, hidden_dim),
            nn.LeakyReLU(0.2),
            nn.Linear(hidden_dim, 1),
            nn.Sigmoid(), # Output probability between 0 and 1.
        )

    def forward(self, x):
        return self.net(x)

```

2.3 GAN Training Loop

In this cell you'll write the training loop for the GAN. Pay close attention to the following steps:

1. Discriminator Training:

- **Real Data:**
Sample a batch of real Swiss Roll points and assign them the label 1.
- **Fake Data:**
Sample random noise, generate fake points with the Generator, and assign them the label 0.
- **Discriminator Loss:**
Compute the Binary Cross Entropy (BCE) loss for both real and fake samples and sum them.

2. Generator Training:

- Generate fake data from random noise.
- Compute the generator loss by passing the fake samples through the discriminator and comparing them to the true label (1) because the generator wants to fool the discriminator.

TODO:

Implement (or study and understand) the loss calculations for both generator and discriminator. Check that: - The discriminator uses correct true labels for real data (1) and fake data (0). - The generator is trained with fake data but with the target label set to 1.

HINTS are provided in the inline comments.

```
[6]: def train_gan(n_steps=10000, batch_size=128, lr=2e-4, latent_dim=2):
    """
    Train the GAN on Swiss Roll data.

    Args:
        n_steps (int): Number of training steps.
        batch_size (int): Batch size for training.
        lr (float): Learning rate.
        latent_dim (int): Dimension of the noise input for the generator.

    Returns:
        generator, discriminator, g_losses, d_losses, device
    """
    device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

    # Initialize networks.
    generator = Generator(latent_dim=latent_dim).to(device)
    discriminator = Discriminator().to(device)

    # Set up optimizers with recommended beta parameters for GAN training.
    g_optimizer = torch.optim.Adam(generator.parameters(), lr=lr, betas=(0.5, 0.
    ↪999))
    d_optimizer = torch.optim.Adam(
        discriminator.parameters(), lr=lr, betas=(0.5, 0.999)
    )

    # Define the binary cross-entropy loss function.
    criterion = nn.BCELoss()

    # Containers for losses.
    g_losses, d_losses = [], []

    from tqdm import tqdm # Ensures that tqdm is imported if not already.

    pbar = tqdm(range(n_steps), desc="Training GAN")
```

```

for step in pbar:
    # ----- Train Discriminator -----
    for _ in range(1):
        # Sample a batch of real data.
        real_data = generate_swiss_roll(batch_size).to(device)
        real_labels = torch.ones(batch_size, 1, device=device) # Real
        ↪ labels are 1.

        # Sample noise from the latent distribution and generate fake data.
        z = torch.randn(batch_size, latent_dim, device=device)
        fake_data = generator(
            z
        ).detach() # Detach to avoid backprop into generator.
        fake_labels = torch.zeros(
            batch_size, 1, device=device
        ) # Fake labels are 0.

        # Zero the gradients for discriminator.
        d_optimizer.zero_grad()

        # Forward pass through the discriminator.
        real_preds = discriminator(real_data)
        fake_preds = discriminator(fake_data)

        # Compute discriminator loss:
        # 1. Use your chosen criterion on (real_preds, real_labels) to
        ↪ calculate the loss for real samples.
        # 2. Do the same for fake samples with (fake_preds, fake_labels).
        # 3. Sum or average these two losses appropriately.
        # Replace the following line with your implementation.

        d_real_loss = criterion(real_preds, real_labels)
        d_fake_loss = criterion(fake_preds, fake_labels)
        d_loss = d_real_loss + d_fake_loss

        d_loss.backward()
        d_optimizer.step()

    # ----- Train Generator -----
    for _ in range(1):
        # Sample noise to generate fake data.
        z = torch.randn(batch_size, latent_dim, device=device)
        g_optimizer.zero_grad()
        fake_data = generator(z)

        # Compute the generator loss:

```

```

        # 1. Pass fake_data through the discriminator.
        # 2. Compare the output with real labels (ones) because the
        ↪ generator's goal is to fool the discriminator.
        # Replace the following line with your implementation.

        g_loss = criterion(fake_preds, real_labels)

        g_loss.backward()
        g_optimizer.step()

    # Record losses.
    g_losses.append(g_loss.item())
    d_losses.append(d_loss.item())

    if step % 100 == 0:
        pbar.set_postfix({"G_Loss": g_loss.item(), "D_Loss": d_loss.item()})

    return generator, discriminator, g_losses, d_losses, device

# Train the GAN.
generator, discriminator, g_losses, d_losses, device = train_gan()

# %% [code]
plt.plot(g_losses, label="Generator Loss")
plt.plot(d_losses, label="Discriminator Loss")
plt.xlabel("Steps")
plt.ylabel("Loss")
plt.title("GAN Training Loss")
plt.legend()
plt.show()

```

Training GAN: 0%| | 0/10000 [00:00<?, ?it/s]

Training GAN: 0%| | 0/10000 [00:00<?, ?it/s]

```

-----
RuntimeError                                Traceback (most recent call last)
Cell In[6], line 100
    96     return generator, discriminator, g_losses, d_losses, device
    99 # Train the GAN.
--> 100 generator, discriminator, g_losses, d_losses, device = train_gan()
    102 # %% [code]
    103 plt.plot(g_losses, label="Generator Loss")

Cell In[6], line 86, in train_gan(n_steps, batch_size, lr, latent_dim)
    79     # Compute the generator loss:
    80     # 1. Pass fake_data through the discriminator.

```

```

81     # 2. Compare the output with real labels (ones) because the
↳generator's goal is to fool the discriminator.
82     # Replace the following line with your implementation.
84     g_loss = criterion(fake_preds, real_labels)
--> 86     g_loss.backward()
87     g_optimizer.step()
89     # Record losses.

```

```

File c:\Users\GinMa\.conda\envs\CS7150\lib\site-packages\torch\_tensor.py:525,
↳in Tensor.backward(self, gradient, retain_graph, create_graph, inputs)
515 if has_torch_function_unary(self):
516     return handle_torch_function(
517         Tensor.backward,
518         (self,),
519         (...)
523         inputs=inputs,
524     )
--> 525 torch.autograd.backward(
526     self, gradient, retain_graph, create_graph, inputs=inputs
527 )

```

```

File c:\Users\GinMa\.conda\envs\CS7150\lib\site-packages\torch\autograd\_init_
↳py:267, in backward(tensors, grad_tensors, retain_graph, create_graph,
↳grad_variables, inputs)
262     retain_graph = create_graph
264 # The reason we repeat the same comment below is that
265 # some Python versions print out the first line of a multi-line functio
266 # calls in the traceback and some print out the last line
--> 267 _engine_run_backward(
268     tensors,
269     grad_tensors_,
270     retain_graph,
271     create_graph,
272     inputs,
273     allow_unreachable=True,
274     accumulate_grad=True,
275 )

```

```

File c:\Users\GinMa\.conda\envs\CS7150\lib\site-packages\torch\autograd\graph.p :
↳744, in _engine_run_backward(t_outputs, *args, **kwargs)
742     unregister_hooks = _register_logging_hooks_on_whole_graph(t_outputs
743 try:
--> 744     return Variable._execution_engine.run_backward( # Calls into the
↳C++ engine to run the backward pass
745         t_outputs, *args, **kwargs
746     ) # Calls into the C++ engine to run the backward pass
747 finally:
748     if attach_logging_hooks:

```

```
RuntimeError: Trying to backward through the graph a second time (or directly
↳access saved tensors after they have already been freed). Saved intermediate
↳values of the graph are freed when you call .backward() or autograd.grad().
↳Specify retain_graph=True if you need to backward through the graph a second
↳time or if you need to access saved tensors after calling backward.
```

2.4 Visualizing GAN Results

Finally, we visualize samples generated from the trained generator alongside real Swiss Roll data.

```
[ ]: def visualize_gan_samples(generator, n_samples=1000):
    """
    Visualize samples from the generator along with real data.

    Args:
        generator: Trained generator model.
        n_samples: Number of samples to generate.
    """
    device = next(generator.parameters()).device

    with torch.no_grad():
        # Generate fake samples.
        z = torch.randn(n_samples, 2, device=device)
        fake_data = generator(z).cpu().numpy()

    # Get real data for comparison.
    real_data = generate_swiss_roll(n_samples).numpy()

    # Create side-by-side plots.
    plt.figure(figsize=(12, 5))

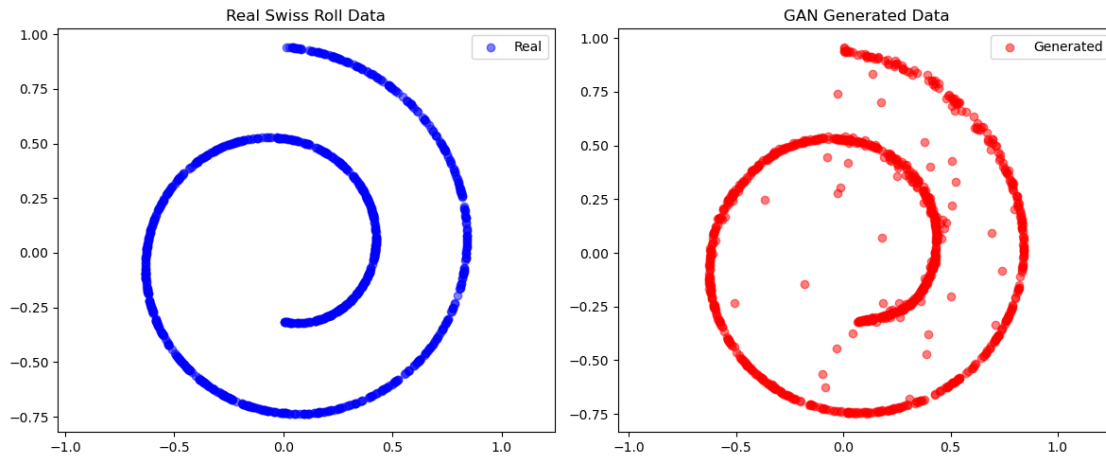
    plt.subplot(1, 2, 1)
    plt.scatter(real_data[:, 0], real_data[:, 1], c="blue", alpha=0.5,
↳label="Real")
    plt.title("Real Swiss Roll Data")
    plt.legend()
    plt.axis("equal")

    plt.subplot(1, 2, 2)
    plt.scatter(fake_data[:, 0], fake_data[:, 1], c="red", alpha=0.5,
↳label="Generated")
    plt.title("GAN Generated Data")
    plt.legend()
    plt.axis("equal")

    plt.tight_layout()
```

```
plt.show()

# Visualize the generated results.
visualize_gan_samples(generator)
```



3 Extra Resource for Reference

3.1 Understanding GANs

Imagine walking into an art studio where something fascinating is happening: there's an artist and an art critic engaged in an unusual competition. The artist is trying to create paintings that look exactly like works from a famous gallery, while the critic's job is to spot the difference between real gallery paintings and the artist's creations. This back-and-forth between artist and critic perfectly captures how GANs (Generative Adversarial Networks) work.

3.2 Core Intuition

In our Swiss Roll example, we can think of it this way: The artist (Generator) is trying to create points that form the distinctive spiral pattern of a Swiss Roll, while the critic (Discriminator) needs to distinguish between points from the real Swiss Roll distribution and the artist's generated points. Through this competition, both networks improve – the artist gets better at creating realistic points, and the critic becomes more discerning.

3.3 The Mathematical Dance

This artistic competition translates into a mathematical game. The Generator (G) and Discriminator (D) play what we call a minimax game, represented by this equation:

$$\min_G \max_D \mathbb{E}_x[\log D(x)] + \mathbb{E}_z[\log(1 - D(G(z)))]$$

Breaking this down in simple terms:

- $D(x)$ is the critic's judgment of real data (should be close to 1)
- $G(z)$ is the artist creating new data from random noise z
- $D(G(z))$ is the critic's judgment of generated data (should be close to 0)
- The Generator tries to minimize this difference
- The Discriminator tries to maximize it

3.3.1 The Latent Space (potential homework? to explore)

The noise input to our Generator isn't just random numbers – it's a structured space where:

- Similar points generate similar outputs
- We can interpolate between points
- Different dimensions might control different features

3.4 Key Papers

For a deeper dive into GANs, consider reading:

- “Generative Adversarial Networks” by Goodfellow et al. (2014)
- “Wasserstein GAN” by Arjovsky et al. (2017)